

# *A Performance Study of a Distributed Database Using TPC-C*

Adil Rahman

*Department of Computer Science*  
University of California, Santa Cruz  
California, United States  
anrahman@ucsc.edu

**Abstract** — Transitioning from centralized to distributed databases has become a popular topic for research in addressing scalability issues. This report seeks to implement a distributed database system that overcomes issues that centralized systems face such as throughput bottlenecks and limited scalability. The initial prototype implemented in Python was created with serial and concurrent locking mechanisms. Throughput for varying number of interconnected nodes ranged from 2,290-11,681 operations/sec and 206-292 tpmC. Minimum latency values for four interconnected nodes with concurrent execution for 1,000 and 10,000 operations were 0.2643 and 2.6296 sec respectively. This establishes a good baseline result for a distributed database design. Implementing the model on different programming language platforms as well as more sophisticated hardware is a potential further research direction.

**Keywords** — *database, distributed, node, serial, serializability, concurrency, lock, scalability, throughput, fault-tolerance, latency*

## I. INTRODUCTION

The need for sophisticated database architectures is rapidly increasing due to the large amounts of data being generated. In addition to this, more requests are being made to that data by the increasing user population demanding database access utilized by various applications.

Many applications currently implement centralized database systems. In certain areas, using a centralized model can be beneficial, as they provide advantages such as data integrity maximization, accuracy, and consistency. Centralized databases have been created via database virtualization techniques and are shown to reduce overheads for data collection in data mining [1]. Other applications that need more protection in terms of security also may choose to utilize centralized databases. System administrators can control and restrict user access more easily when all data is stored in one physical location, as seen in some applications such as smart cards [2].

However, major problems that these architectures face include meeting user requests for specific data segments [3]. If too many users are requesting various data from one centralized database, it can create a bottleneck that leads to suboptimal efficiency in processing and managing transactions. Centralized approaches also face problems in terms of data transmission of files or resources. These approaches are not realistic to implement in terms of scalability of applications that have a large number of end users/systems. This can easily reach total data transmissions on the GB scale per hour [4]. In addition to this, centralized systems are more prone to downtime which can impact users

if system failures or malicious attacks occur. Although data integrity is a defending point to this model, we must consider the case that a centralized system's security is not up to par. The integrity of the entire database can potentially be compromised. Delta Airlines was subject to this, causing numerous flights to be canceled and impacted the business severely [5].

Distributed databases differ from their centralized counterparts in that there are now a collection of databases that are located at different sites, rather than one physical location. A distributed database management system (DDMS) must be present that can take care of query processing and structured data organization on top of the functionality that a standard database management system would take care of [6]. When users send transactions to be executed, the DDMS ensures that the transactions involving shared data between uses do not conflict with each other. If data is shared among the transactions, a scheduler is responsible for ensuring the database is still correct after concurrent accesses are performed by multiple transactions. This ensures concurrency control for both data integrity and serializability.

There are different ways that a DDMS can enforce concurrency control. One successful approach is a locking mechanism. If two transactions conflict, the scheduler can make one transaction wait for the other transaction to perform its operations first before performing its own. Read and write locks are placed on a datum. If a write lock is present, no other transaction can obtain another write or read lock on the same datum. A read lock can be obtained if other read locks are on the datum [7].

One example of a concurrency control method that guarantees serializability is two-phase locking (2PL) [11]. This protocol ensures that two transactions cannot hold conflicting locks on the same data items at the same time. 2PL consists of the growing phase and the shrinking phase [10]. The growing phase consists of acquiring all locks and not releasing any locks. A transaction performs once it obtains all necessary locks. The shrinking phase then consists of releasing all locks and not acquiring any new locks.

An issue that can occur in concurrency control is when two or more transactions are waiting for each other to release locks of shared data between the transactions in the growing phase. This is referred to as a deadlock. One variant of 2PL that addresses this issue is conservative two-phase locking (C2PL) [12]. C2PL prevents deadlocks by ensuring

that all locks that a given transaction needs must be obtained first before the transaction performs.

The distributed database architecture addresses bottleneck issues. Users can request information from specific servers that have the data that they need, without having to go through a central authority. Additionally, total data transmissions are now reduced and distributed across multiple servers which can free up the potential load on a user's end. Lastly, distributed systems are robust and fault-tolerant when system failures or malicious attacks occur. Distributed databases have shown to provide increased reliability and performance as seen in the relational model of distributed Ingres [8].

The Paxos algorithm is the standard tool for implementing a fault-tolerant distributed system [13]. The Paxos protocol ensures that consensus is achieved among many processes in which a single value that has been proposed is chosen. Three types of roles are established for each process, which consists of proposers, acceptors, and learners. Proposer processes will propose values, while many acceptor proposers will choose a single proposed value. Learner processes will then learn a value if one has been chosen. Other processes will only learn that a value has been chosen when the single value is chosen. The original Paxos algorithm allows for  $f$  malicious processes so long as there are  $2f + 1$  non-faulty processes.

The primary objective of this report is to develop a distributed database system that will use a C2PL locking mechanism to enforce concurrency control for transactions. The prototype will be implemented in Python.

## II. IMPLEMENTATION

The application implements a communication interface that supports sending messages between any two servers. Each server is represented as a physical node that can send and receive messages between other nodes with an established connection. Each node is responsible for a specific shard of the database. If a node needs to access a specific data element that is not under its domain, it will communicate with the node that has that data under its domain and requests that node to perform the necessary transaction.

For the proposed application, each transaction consists of a single read or write operation, although transactions can consist of larger amounts of operations. Each transaction receives a lock corresponding to a datum before it performs its transaction. If a datum is currently being worked on by transaction  $T_1$ , then any other transactions trying to read or write it must wait until  $T_1$  has committed and completed its operation. The following sections describe the process in more detail.

### A. Nodes

Let  $n$  represent the number of interconnected nodes. Each node performs both server and client roles. Fig. 1 shows a diagram of a network of interconnected nodes where each node serves as both a server and a client. Each node listens for potential connections from other nodes using a socket bound to a specific port. Every node is

capable of connecting to other nodes and maintains a list of node addresses and ports to which it is currently connected to. A node remains connected to any number of nodes after the initial connection has been established. Each node also keeps track of its own partition of the data which is determined by a static mapping mechanism.

To execute a transaction, the client sends a message to the node that is holding the data that is being accessed as part of the transaction. The server component of the node responsible for fulfilling the request will then receive that transaction. The server component of this node has a transaction manager that will process and perform the transaction.




Fig. 1 Interconnected node network with server and client roles for each node where  $n = 3$ .

The transaction manager achieves this by analyzing the string encoded message received and then looks for keywords that indicate a transaction type that needs to be performed. If the transaction manager does not recognize a particular transaction type from the received message, it will indicate to the client node that the message was unrecognizable, returns false, and nothing was done to its partition of the database. If the transaction is valid, then the server performs it and sends a message respective to the transaction type back to the client node that originally had issued the transaction request.

### B. Physical distribution of data

The database is simulated by utilizing a key-value store in which the number of data items each node is responsible for is a parameter set by the user. Each node initiates its domain which is tied to the data that the node is responsible for. Each time a new node connects, it determines the domain of data items it is responsible for based on the existing node's domains.

Let  $x_i$  represent the maximum threshold that a particular node  $N_i$  is responsible for. If  $n = 1$ , then  $x_1 = 100$  and  $N_1$  will have a range of data items  $\{0, 99\}$ . If  $n = 2$ , then  $x_1 = 100$ , and  $x_2 = x_1 + 100 = 200$  and  $N_1$  will still have a domain of data items  $\{0, 99\}$  while  $N_2$  will have a domain

of data items  $\{100, 199\}$ . Thus when  $n = k$ , then  $x_k = x_{k-1} + 100$  and  $N_k$  will have a domain of data items  $\{x_{k-1}, x_k - 1\}$ . In this way, data responsibility is distributed across any number of  $k$  nodes. Fig. 2 shows an illustrative diagram of how each node domain/responsibility appears within the model.




Fig. 2 Domain responsibility of  $n = 3$  nodes and potential additional nodes.




Fig. 3 Example of  $n = 2$  nodes sending and receiving messages to each other.  $N_1$  who is blue sends to red and  $N_2$  who is red receives from blue.  $N_2$  who is red sends to blue and  $N_1$  who is blue receives from red.

### C. Sending and receiving messages

Each node is capable of sending and receiving messages to other nodes. If  $n = 1$ , a node does not send or receive messages to other nodes, as it is the only node present. The single node will be able to perform operations on its partition of the database freely. This node's partition of the database will be equivalent to the entire current database.

If  $n > 1$ , then a node will be able to send and receive messages to any node that it is connected to. Each node executes two threads, where the first thread is responsible for sending messages to other connected nodes, while the second thread is responsible for receiving messages from other connected nodes. Fig. 3 shows a diagram of two nodes sending and receiving messages via threads. Both of these threads loop continuously so that if at any time any node including itself tries to perform a transaction, all nodes can attempt to perform it. If a client component of a node requests for a transaction to be performed, it will perform all of the operations in the

transaction that affect data in its domain. It will then send any operations for data outside of its domain to the respective nodes that have the data in their domain. These nodes will then perform the operations and send a message back to the node who initiated the transaction to inform it that the operations were either committed or not.

### D. Locking Mechanism

Before any transaction or operation is performed, a specific thread must obtain a lock on a particular data item. A list of locks is created corresponding to the number of data items that are present across  $n$  interconnected nodes. If there are  $n$  nodes, then there are also  $n$  number of partitions of the database. For instance, if each partition holds 100 data items, the maximum number of locks present at any time will be equal to  $100n$ . Each lock and data item shares the same key identity. If a transaction needs to perform an operation on a specific data item at a key and corresponding value, then it will have to also obtain the lock corresponding to that key. Since the key identities for both lock and values are shared, the key-value store and lock list exhibit a parallel characteristic. Fig. 4 shows the overarching parallel structure of each node's partition of the data and the corresponding locks.

Once a read or write operation needs to be performed, a thread will be created in the transaction that will obtain a lock for the data item it needs to read or write. The thread will perform the operation and then release the lock. The scenario in which two or more transactions  $T_1$  and  $T_2$  are attempting to perform a read or write operation on the same data item is possible. If  $T_1$  obtains the lock first, then  $T_2$  must wait for  $T_1$  to commit its operation before  $T_2$  can perform any operations on the specific datum.  $T_2$  will indefinitely attempt to obtain the lock that  $T_1$  currently has. Once  $T_1$  has finished, it will release its lock, allowing  $T_2$  to now obtain the lock and perform its operation. Since each transaction consists of one read or write operation, this emulates a C2PL mechanism and prevents the possibility of deadlocks.

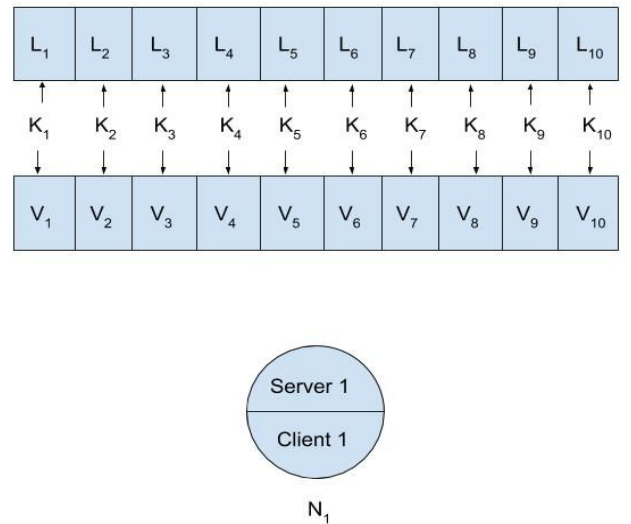


Fig. 4 Example of one node's parallel structure of values and lock sharing the same key. Example only shows 10 locks and values, however, there exists the same structure for all 100 data items in a node's domain.

### E. Read and Write Operations

The transactions contain read or write operations that will access data in multiple distributed partitions. A read operation takes a key as input and returns the corresponding value currently in the key-value store. A write operation takes a key and a value that overwrites the current value currently in the key-value store. When a node performs a read operation that is outside its domain, the read operation request will be sent to the node who has the data item in its domain, perform the read operation, and send the value back to the node who requested it. When a node performs a write operation that is outside its domain, the write operation request will be sent to the node who has the data item in their domain. It will then perform the write operation and send a message that the write was successful to the requester node.

### F. Benchmark Evaluation

There is a need to test the distributed database prototype on a workload representing a more complex application environment in addition to the basic read and write operations mentioned above. The TPC Benchmark™ C (TPC-C) is characterized as an emulated online transactional processing (OLTP) workload that is considered the standard for evaluating database performance. TPC-C consists of five transaction profiles that consist of a mix of read-only and update-intensive transactions. These profiles are performed on a collection of nine relational tables.

The TPC-C benchmark represents a wholesale supplier company where the TPC-C database maintains information for nine separate and individual tables. These tables contain various information such as unique IDs, variable text, fixed text, date and time, and signed and unsigned numerics. Each terminal server will represent a warehouse, where each warehouse is responsible for ten districts under it. Each of those districts serves three thousand customers. Each server has its own tables to keep track of their items and stock, as well as history for new orders. In our server client model, each server node is also acting as a client customer, where a given transaction profile can be executed. When a transaction profile is requested to be performed by the client portion, the server portion processes and performs the transaction and ultimately returns the result to the client portion.

The TPC-C benchmark consists of five transaction profiles. The *New-Order* transaction represents a medium intensity read and write transaction where a customer completes a new order for some items. The *Payment* transaction represents a light intensity read and write transaction where the balance information will be updated for a given customer based on a set of purchased items. The *Order-Status* transaction represents a medium read-only transaction in which a customer can query the status of their last order. The *Delivery* transaction represents a batch of read write transactions where new orders that are within a queue are processed. The *Stock-Level* transaction represents a heavy read-only transaction where the stock levels that are below a specified threshold are identified.

In our model, each node maps to a single warehouse. The transaction input parameters randomly choose whether to request a resource in the home warehouse or a remote warehouse based on a non-uniform probability distribution that is biased towards the home warehouse. This requires that a given server warehouse request data that is located under another server's domain. The performance metric that is reported through this workload is the measurement of the number of orders processed per minute, which is expressed as transactions-per-minute-C (tpmC).

## III. EVALUATION STUDY

The implementation above was tested for throughput, latency, scalability, and fault-tolerance. Tests were performed for three different scenarios. The first scenario runs the implementation without a locking mechanism implemented and is represented as read no locking (RNL) and write no locking (WNL). The second scenario implements a locking mechanism in which each operation is forced to act serially and is represented as read test serial (RTS) and write test serial (WTS). The third scenario implements a locking mechanism that allows operations to act concurrently and is represented as read test concurrent (RTC) and write test concurrent (WTC). All scenarios were tested for read-only and write-only transactions. Lastly, the TPC-C benchmark implementation is run on a varying number of  $n$  interconnected nodes and the throughput of TPC-C transaction profiles is obtained. The following sections explore the results.

### A. Throughput

Throughput was measured to evaluate the number of read or write operations that can be performed in one second for all scenarios with  $n = 1$  to  $n = 4$  interconnected nodes. The difference in the number of operations that are performed in one second at  $n = 1$  is negligible. This is because a single node does not need to send any messages to other nodes that take up more time. However, as more nodes are interconnected, the number of operations that fall within the requesting node's domain decreases. As more and more nodes are involved in the reading and writing of various data items across several nodes domains, the number of operations per second decreases as more communication between other interconnected nodes is necessary for completing a transaction. It should be noted that this model does not persist writes to a hard disk. Due to this, read and write operations take a similar amount of time. Table I reflects this property as the number of reads and writes is similar, as well as the decreasing nature of throughput as  $n$  interconnected nodes increase.

TABLE I. Throughput

# of Nodes	Throughput (operations / sec)					
	RNL	RTS	RTC	WNL	WTS	WTC
$n = 1$	11,681	11,469	11,486	11,623	11,443	11,527
$n = 2$	3,741	3,437	3,631	4,951	4,492	4,909
$n = 3$	3,201	2,901	3,032	4,289	3,755	4,134
$n = 4$	2,885	2,290	2,725	3,794	2,760	3,761

## B. Latency

Latency was measured based on 1,000 and 10,000 operations for all scenarios with  $n = 1$  to  $n = 4$  interconnected nodes. In the case where only one node is present, the difference between no locking, serial execution, and concurrent execution of operations is small. However, as the amount of interconnected nodes increases, the number of operations increases, the difference in latency is much more apparent, as seen in Tables II - V.

All scenarios that do not implement a locking system (RNL and WNL) take the least amount of time to perform. Since no lock needs to be retrieved per data item, all threads doing read and write operations are free to perform their tasks without waiting for other threads that could be working on the same item.

All scenarios that implement a concurrent execution take approximately 1% more time than the scenario of a no locking mechanism. In the concurrent execution, there may be instances where a thread  $th_2$  is waiting to retrieve a lock that a pre-existing thread  $th_1$  performing an operation on the same key currently has. This will increase the latency slightly as  $th_2$  must wait for  $th_1$  to release the lock so that it can acquire it. This scenario could occur where any  $p$  number of threads could all be trying to perform an operation on the same data item, in which the concurrency model will slow down.

Lastly in the serial execution, there is only one lock present and must be acquired before doing any operations whatsoever. This forces all threads to have to wait for the transaction before it to complete its operation before attempting to do their own, which slows down latency. Tables II - V demonstrates the varying latencies of all scenarios as well as varying numbers of interconnected nodes.

TABLE II. Latency,  $n = 1$

Number of Operations	Latency (sec)					
	RNL	RTS	RTC	WNL	WTS	WTC
1,000	0.0922	0.0935	0.0924	0.0891	0.0946	0.0943
10,000	0.8043	0.8131	0.8123	0.8163	0.8327	0.8317

TABLE III. Latency,  $n = 2$

Number of Operations	Latency (sec)					
	RNL	RTS	RTC	WNL	WTS	WTC
1,000	0.2659	0.2812	0.2672	0.2026	0.2261	0.2033
10,000	2.6055	2.7567	2.6635	1.9426	2.1704	2.0080

TABLE IV. Latency,  $n = 3$

Number of Operations	Latency (sec)					
	RNL	RTS	RTC	WNL	WTS	WTC
1,000	0.3031	0.3488	0.3171	0.2237	0.2626	0.2362
10,000	3.1392	4.7744	3.1596	2.2272	3.7824	2.3291

TABLE V. Latency,  $n = 4$

Number of Operations	Latency (sec)					
	RNL	RTS	RTC	WNL	WTS	WTC
1,000	0.3489	0.3658	0.3510	0.2637	0.2744	0.2643
10,000	3.5833	4.9337	3.6785	2.5396	3.9570	2.6296

## C. Scalability

In this implementation, each node was responsible for only 100 data items. However, as seen in sections A and B, with an increasing number of nodes, fewer operations can be performed because it requires more time to perform operations. This is due to read and write operations having to be sent to other nodes more frequently as the number of  $n$  interconnected nodes increases. If the domains of each individual node are small, then the likelihood a node will have to send an operation to another node is much higher.

Thus one way to increase the scalability of this model is to increase the domain size of each individual node. In a real application setting, this number can be much higher than 100 and can be set to any number that is desired. By increasing the domain size, the likelihood that a server requests operations to be done from another server is lower since there is a higher chance the operation it needs to do is already in its domain. Depending on the application, data can be strategically placed in different servers so that users on the application level can retrieve information relevant to them much faster without having to rely on retrieval from other servers.

One example of this is a bank account database. If a user resides in Santa Cruz, California, then the user's account information is stored in a server that is local to that site, either in Santa Cruz or closeby. The user would be able to retrieve and write information very quickly in regards to their account without any other information needing to be accessed.

## D. Fault-Tolerance

The current implementation has nodes keeping track of which other nodes they are connected to. If a node is to run into system failures and then become disconnected from other nodes, each node that is still connected will still be able to be functional. This includes being able to continue to perform read and write operations for their domains and request nodes that are still connected to perform operations on their respective domains.

However, one side effect of a node failing is that its entire partition of the data will be completely lost. The current implementation does not persist writes to any hard disk. If the randomization of input parameters happens to include portions of data that were within the lost partition, the nodes will also fail. One way to address this is to back up the data for each node in a separate physical entity. A node would then try to send a message to the fallen node. If after some time frame there is no response from the fallen node, the node that needs the data would instead access the separate physical entity that has the data. In addition to this, when a fallen node is made available again, it can retrieve all of the information for data in its domain that it previously had.

Another side effect of direct partitioning of data is the scenario of a malicious attacker (MA). If an MA can edit information in the database wrongfully or take control of the node, there is no way to guarantee that the information being read from data in this node is accurate. A solution to this would be to partially replicate data across different nodes. Instead of one node being only responsible for its domain, it would have partially overlapping domains with other nodes. If a malicious node were to be present, a consensus-based decision-making system could be put in place like that of Paxos [13]. The data the malicious node is presenting is compared to the same data that is present in overlapping portions in the other nodes. Then a consensus vote can be implemented to send the correct data where the information sent will be the value that is consistent and stored in the majority of nodes. In this scenario, there must be enough non-malicious nodes to outvote any amount of malicious nodes.

#### E. TPC-C

The tpmC was measured with  $n = 1$  to  $n = 4$  interconnected nodes. Unlike the throughput reported in part A, the tpmC does not change as much as the number of interconnected nodes increases. This is due to the individual nature of the transaction profiles in whether they are completed at home or remotely. Any transaction profiles that require data to be requested from a remote warehouse will add additional time required to complete the transaction. The amount of time spent sending requests for messages and receiving the data back decreases the amount of available time to perform additional transactions.

The Order-Status transaction profile is always processed within the home warehouse and, as a result, does not cause any change in time required to process due to sending data requests from other warehouses. The Delivery transaction profile is a deferred queue that also is processed based on the home warehouse. These two transaction profiles do not affect the tpmC throughput as the number of nodes increases.

The New-Order transaction profile only has a 1% chance that a remote warehouse ID will be generated as input. In most cases, this transaction profile will be completely performed on the partition of the data of the server that called it and not be requested elsewhere. The Payment transaction profile has a 15% chance that a remote warehouse needs to process the payment of the customer. Lastly, the Stock-Level transaction will occur for a randomly chosen warehouse and may cause some added latencies in retrieving the data.

Table VI shows the tpmC for  $n = 1$  to  $n = 4$  nodes. In the case, where there is only one node present, the business throughput is seen at its highest at 292 tpmC. The difference in tpmC is minimal, however, once there exists any number of interconnected nodes greater than one. When a transaction that requires remote data is performed, it will at most have to send a single message to one other server and then retrieve data from that same server. This implies that, regardless of the number of nodes, there will only be a slow down in terms of business throughput as long as the number of nodes is greater than one. This can be seen as the

business throughput for  $n = 2$  to  $n = 4$  nodes is practically the same.

However, as the number of interconnected nodes increases, the amount of time it takes to search for the proper node to request data from will increase. Another time penalty will incur once the node that has gathered the data needs to send the data back to the node that requested the data. A gradual decrease in business throughput is expected as the search space of connected nodes increases.

TABLE VI. TPC-C Results

# of Nodes	Business Throughput (tpmC)
$n = 1$	292
$n = 2$	212
$n = 3$	211
$n = 4$	206

#### IV. CONCLUSION

The initial prototype of a lock mechanism driven distributed database network has been completed. The prototype implementation solves bottleneck issues by partitioning data locally into servers that would cause latency issues in a centralized model. The prototype is fully functional and able to continue running if network nodes fail, unlike a centralized model that only consists of a single node.

Certain limitations may have been imposed by implementing the prototype in Python, which may sacrifice speed for its dynamic typing and interpreted language nature. Transitioning this model to C or Cython may provide performance increases due to these languages not relying on object-oriented programming. In addition to this, future scalability tests will need to be performed on more sophisticated hardware such as a UNIX cluster. Future extensions to the current prototype of this model and further research can be performed to further optimize the communication processes and transactions.

#### ACKNOWLEDGMENT

The author would like to acknowledge the support of Vishal Chakraborty for his help on portions of the programming.

#### REFERENCES

- [1] Y. Wada, Y. Watanabe, K. Syoubu, J. Sawamoto and T. Katoh, "Virtual Database Technology for Distributed Database," *2010 IEEE 24th International Conference on Advanced Information Networking and Applications Workshops*, Perth, WA, 2010, pp. 214-219.
- [2] M. Mohandes, "A smart card management and application system," *2010 IEEE International Conference on Progress in Informatics and Computing*, Shanghai, 2010, pp. 1220-1225.
- [3] M. Y. Jung and J. W. Jang, "Data management and searching system and method to provide increased security for IoT platform," *2017 International Conference on Information and Communication Technology Convergence (ICTC)*, Jeju, 2017, pp. 873-878.
- [4] M. B. Moss, "Comparing Centralized and Distributed Approaches for Operational Impact Analysis in Enterprise Systems," *2007 IEEE*

*International Conference on Granular Computing (GRC 2007)*, Fremont, CA, 2007, pp. 765-765.

- [5] cnn.com, 'Computer glitch grounds Delta flights', 2004. [Online]. Available: <http://www.cnn.com/2004/TRAVEL/05/01/delta.delays/>. [Accessed: 11- Mar - 2019].
- [6] M. T. Ozsu and P. Valduriez, "Distributed database systems: where are we now?," in *Computer*, vol. 24, no. 8, pp. 68-78, Aug. 1991.
- [7] B. Bhargava, "Concurrency control in database systems," in *IEEE Transactions on Knowledge and Data Engineering*, vol. 11, no. 1, pp. 3-16, Jan.-Feb. 1999.
- [8] Stankovic, "A Perspective on Distributed Computer Systems," in *IEEE Transactions on Computers*, vol. C-33, no. 12, pp. 1102-1115, Dec. 1984.
- [9] Transaction Processing Performance Council, "TPC Benchmark C, Standard Specification, Revision 5.11.0", Feb. 2010.
- [10] K. P. Eswaran, J. N. Gray, R. A. Loric, and I. L. Traiger. "The notions of consistence and predicate locks in a database system." *Commun. ACM*, vol. 19, no. 11, pp. 624-633, Nov. 1976.
- [11] P. Leu and B. Bhargava, "Clarification of Two Phase Locking in Concurrent Transaction Processing," *IEEE Trans. Software Eng.*, vol. 14, no. 1, pp. 120-123, 1988
- [12] N.B. Al-Jumah, H.S. Hassanein, and M. El-Sharkawi, "Implementation and modeling of two-phase locking concurrency control - a performance study," *Information and Software Technology*, vol 42, Issue 4, pp. 257-273, 2000.
- [13] L. Lamport, "Paxos Made Simple," *ACM SIGACT News*, 32(4):51-58, 2001.